Vedaang Chopra

vedaangchopra@gatech.edu — +1 (404) 740-9905 — Atlanta, GA — LinkedIn — GitHub

---

Dear Scale AI Hiring Team,

I am applying for the Frontier Agents Engineer (Applied AI) position. My work sits exactly at the intersection this role demands: building agentic AI systems that survive contact with production, and building the evaluation harnesses that prove they work.

Production agents, not prototypes. At Fortinet's AIOps R&D, I led an agentic RAG diagnostics system from hackathon prototype to production deployment: an LLM that plans multi-step tool calls over live network telemetry, invokes APIs, and verifies hypotheses through structured function calling to produce autonomous root-cause analysis – cutting mean resolution time by ∼70%. I scaled OpenSearch pipelines from 50 to 2,000 events/sec along the way, so I have felt the reliability, latency, and observability constraints your enterprise customers operate under.

Multi-agent architecture with real orchestration. My research spans three systems: ATHENA, a five-agent supervisor-worker system with reflection-based quality critique and dynamic plan modification; SHASTRA (in progress), a framework for retrieving and adapting prior agent workflows for new long-horizon tasks under cost/latency budgets; and a closed-loop CAD code-generation pipeline where an LLM's output is executed, verified geometrically, and repaired iteratively – agents grounded by verifiable execution feedback, not single-pass generation.

Evaluation as a first-class discipline. I maintain a 202-test evaluation harness with golden datasets, regression suites, and LLM-as-a-Judge visual scoring that gates every change to my current lab's 16-module pipeline. I have run controlled ablations – removing my ATHENA system's research agent cut fact coverage by 27% in human-rated evaluation (n=5 raters) – because "own the full experimentation lifecycle" is not a slogan; it is how I already work.

What draws me to Scale is the portfolio breadth – solving many different AI problems across industries rather than one narrow slice – and the mandate to pair frontier-model fluency with deterministic engineering. That pairing matches my track record across Python, Go, distributed systems (Azure-deployed), and modern agent tooling.

I would welcome the chance to discuss how my experience building trustworthy agentic systems can contribute to Scale's enterprise deployments.

Sincerely,
Vedaang Chopra